

Frozen ASR-pretrained representations are strong under mild noise but collapse at 5 dB

Train128/dev32 Mini LibriSpeech protocol; all local heads are trained only on clean speech.

